



POLITECNICO
MILANO 1863

POLITECNICO MILANO 1863
NECST
laboratory

MonteCarlo Benchmarking

Andrea Paparella - 10701904

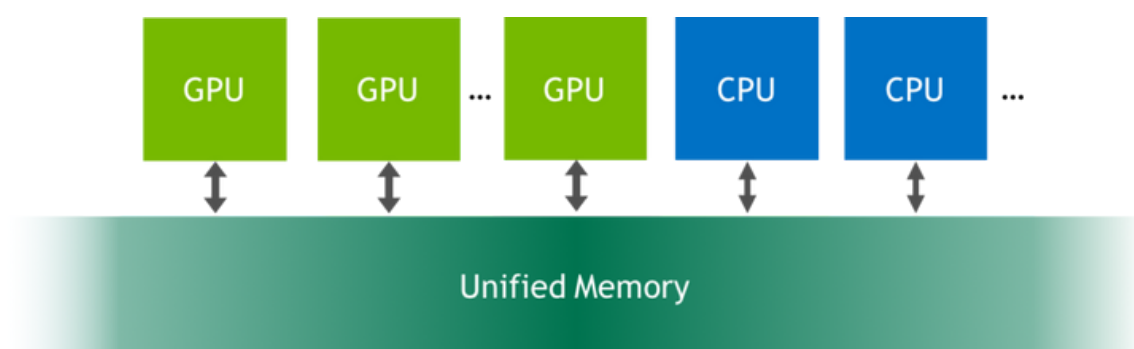
Andrea Piras - 10725972

Professor: Marco Domenico Santambrogio

Advisor: Ian Di Dio Lavore

The Problem

Investigating the performance impact of NVIDIA's Unified Memory paradigm.



Traditional

Allocate memory on CPU and GPU and manually copy data between them.

Unified Memory

Single memory space that is accessible by both the CPU and GPU.

```
checkCudaErrors(cudaMalloc(&plan->d_CallValue, sizeof(__TOptionValue)*(plan->optionCount)));  
checkCudaErrors(cudaMallocHost(&plan->h_OptionData, sizeof(__TOptionData)*(plan->optionCount)));  
//Allocate internal device memory  
checkCudaErrors(cudaMallocHost(&plan->h_CallValue, sizeof(__TOptionValue)*(plan->optionCount)));
```

Original Implementation

From the Tartan benchmark suite
written in C++

- ◆ 2 Modes: Threaded & Streamed
- ◆ 2 Input Scalings: Weak & Strong
- ◆ Allocation: cudaMalloc & cudaMallocHost
- ◆ Data Migration: cudaMemcpyAsync

```
checkCudaErrors(cudaMallocManaged(&plan->um_OptionData, sizeof(__TOptionData) * (plan->optionCount))
checkCudaErrors(cudaMallocManaged(&plan->um_CallValue, sizeof(__TOptionValue) * (plan->optionCount))
// Allocate internal device memory
// Allocate states for pseudo random number generators
```

Unified Memory Implementation

From the original implementation

- ◆ From cudaMalloc & cudaMallocHost to cudaMallocManaged
- ◆ Removal of cudaMemcpyAsync
- ◆ Optimizations through different versions

The Methodology



Performance Metrics

Initialization & Execution Time
Page Faults



Data Collection

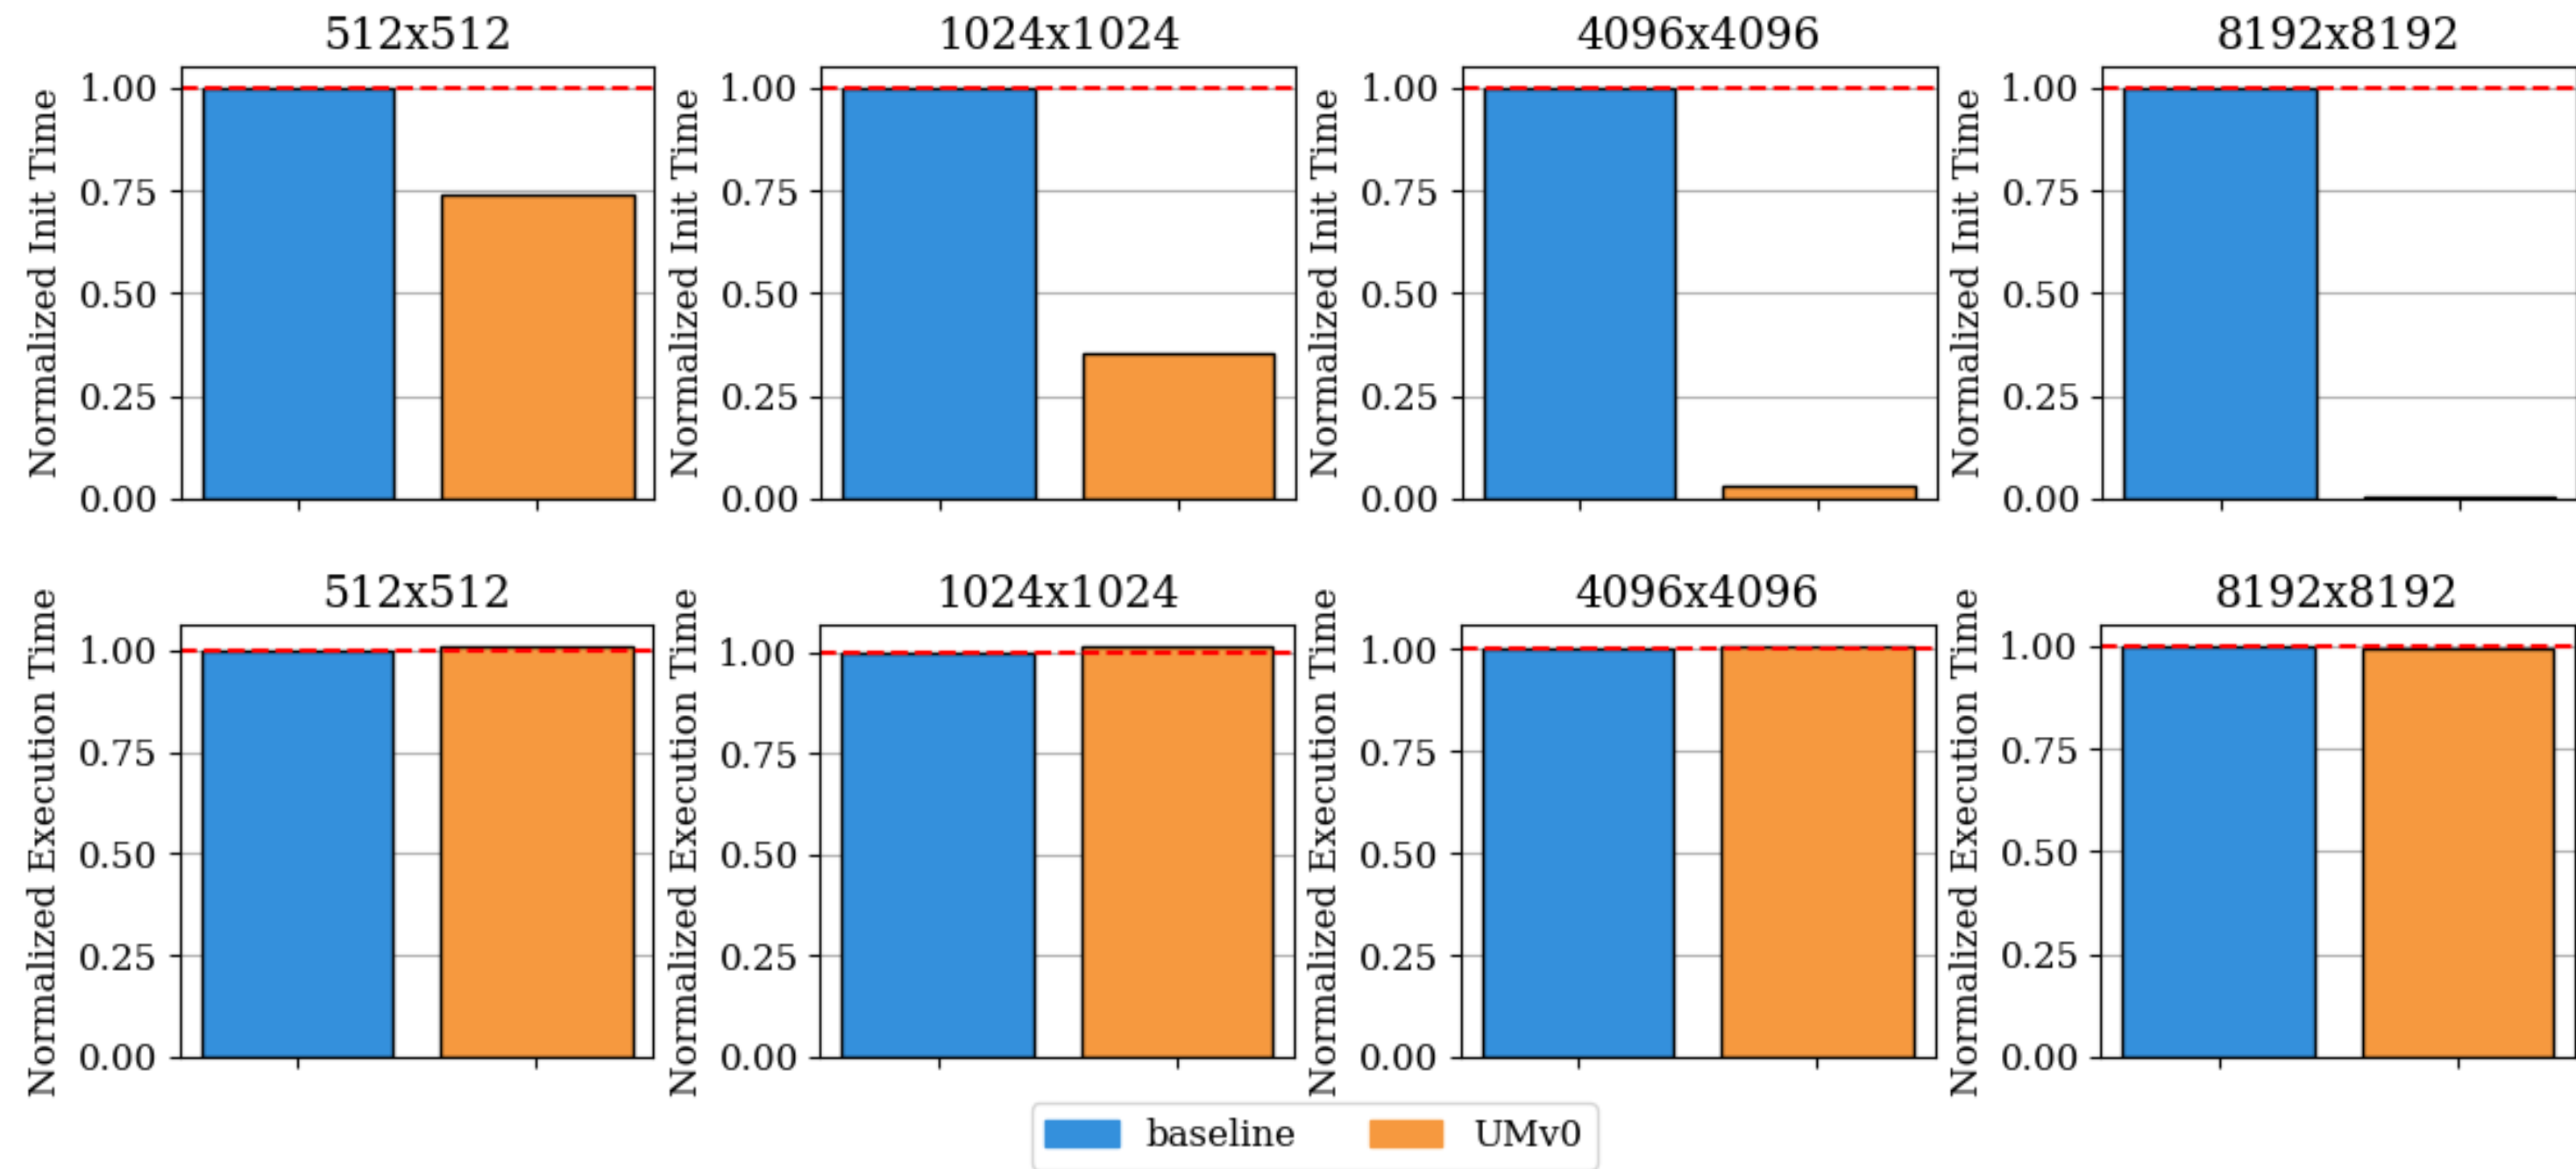
Nvprof & NVIDIA Nsight



Testing

Different Input Sizes
Different GPUs Number
Average Time of 5-10 runs

Results - Unified Memory



Streamed Strong Scaling

Optimizing through prefetching

Version 1

`cudaMemPrefetchAsync`

Proactively allocate memory on a specific location before it is actually accessed

Intuition

Migrate data to device requesting it to avoid page faults caused by attempts to retrieve data that is not present on memory

Optimizing through advising

Version 2

cudaMemAdvise

Advise the Unified Memory subsystem about the memory usage pattern

- `cudaMemAdviseSetPreferredLocation`
- `cudaMemAdviseSetReadMostly`
- `cudaMemAdviseSetAccessedBy`

Intuition

Reduce page faults and memory migrations through automatic optimization guided by the hinted policies

Combining optimizations together

Version 3


Objective

Leverage the optimal features of both versions

Intuition

Prefetching eliminates page faults but can be expensive: replace prefetches with advices to mitigate the occurrence of page faults and enhance performance

Results - Initialization Time



Streamed Strong Scaling



230x

Best SpeedUp



100%


Page Fault
Reduction



% 0.03

Overall Occupancy

Results - Execution Time



Streamed Strong Scaling



<2.5%
Execution Time
Variation



100%
Page Fault
Reduction



% 99.97
Overall Occupancy

```
monteCarloOneBlockPerOptionKernelFunction = cu.invokeMember("bindkernel", path_to_binded_kernels,  
    "cxx MonteCarloOneBlockPerOption( " +  
        "d: out pointer uint32, " +  
        "v: out pointer uint32, " +
```

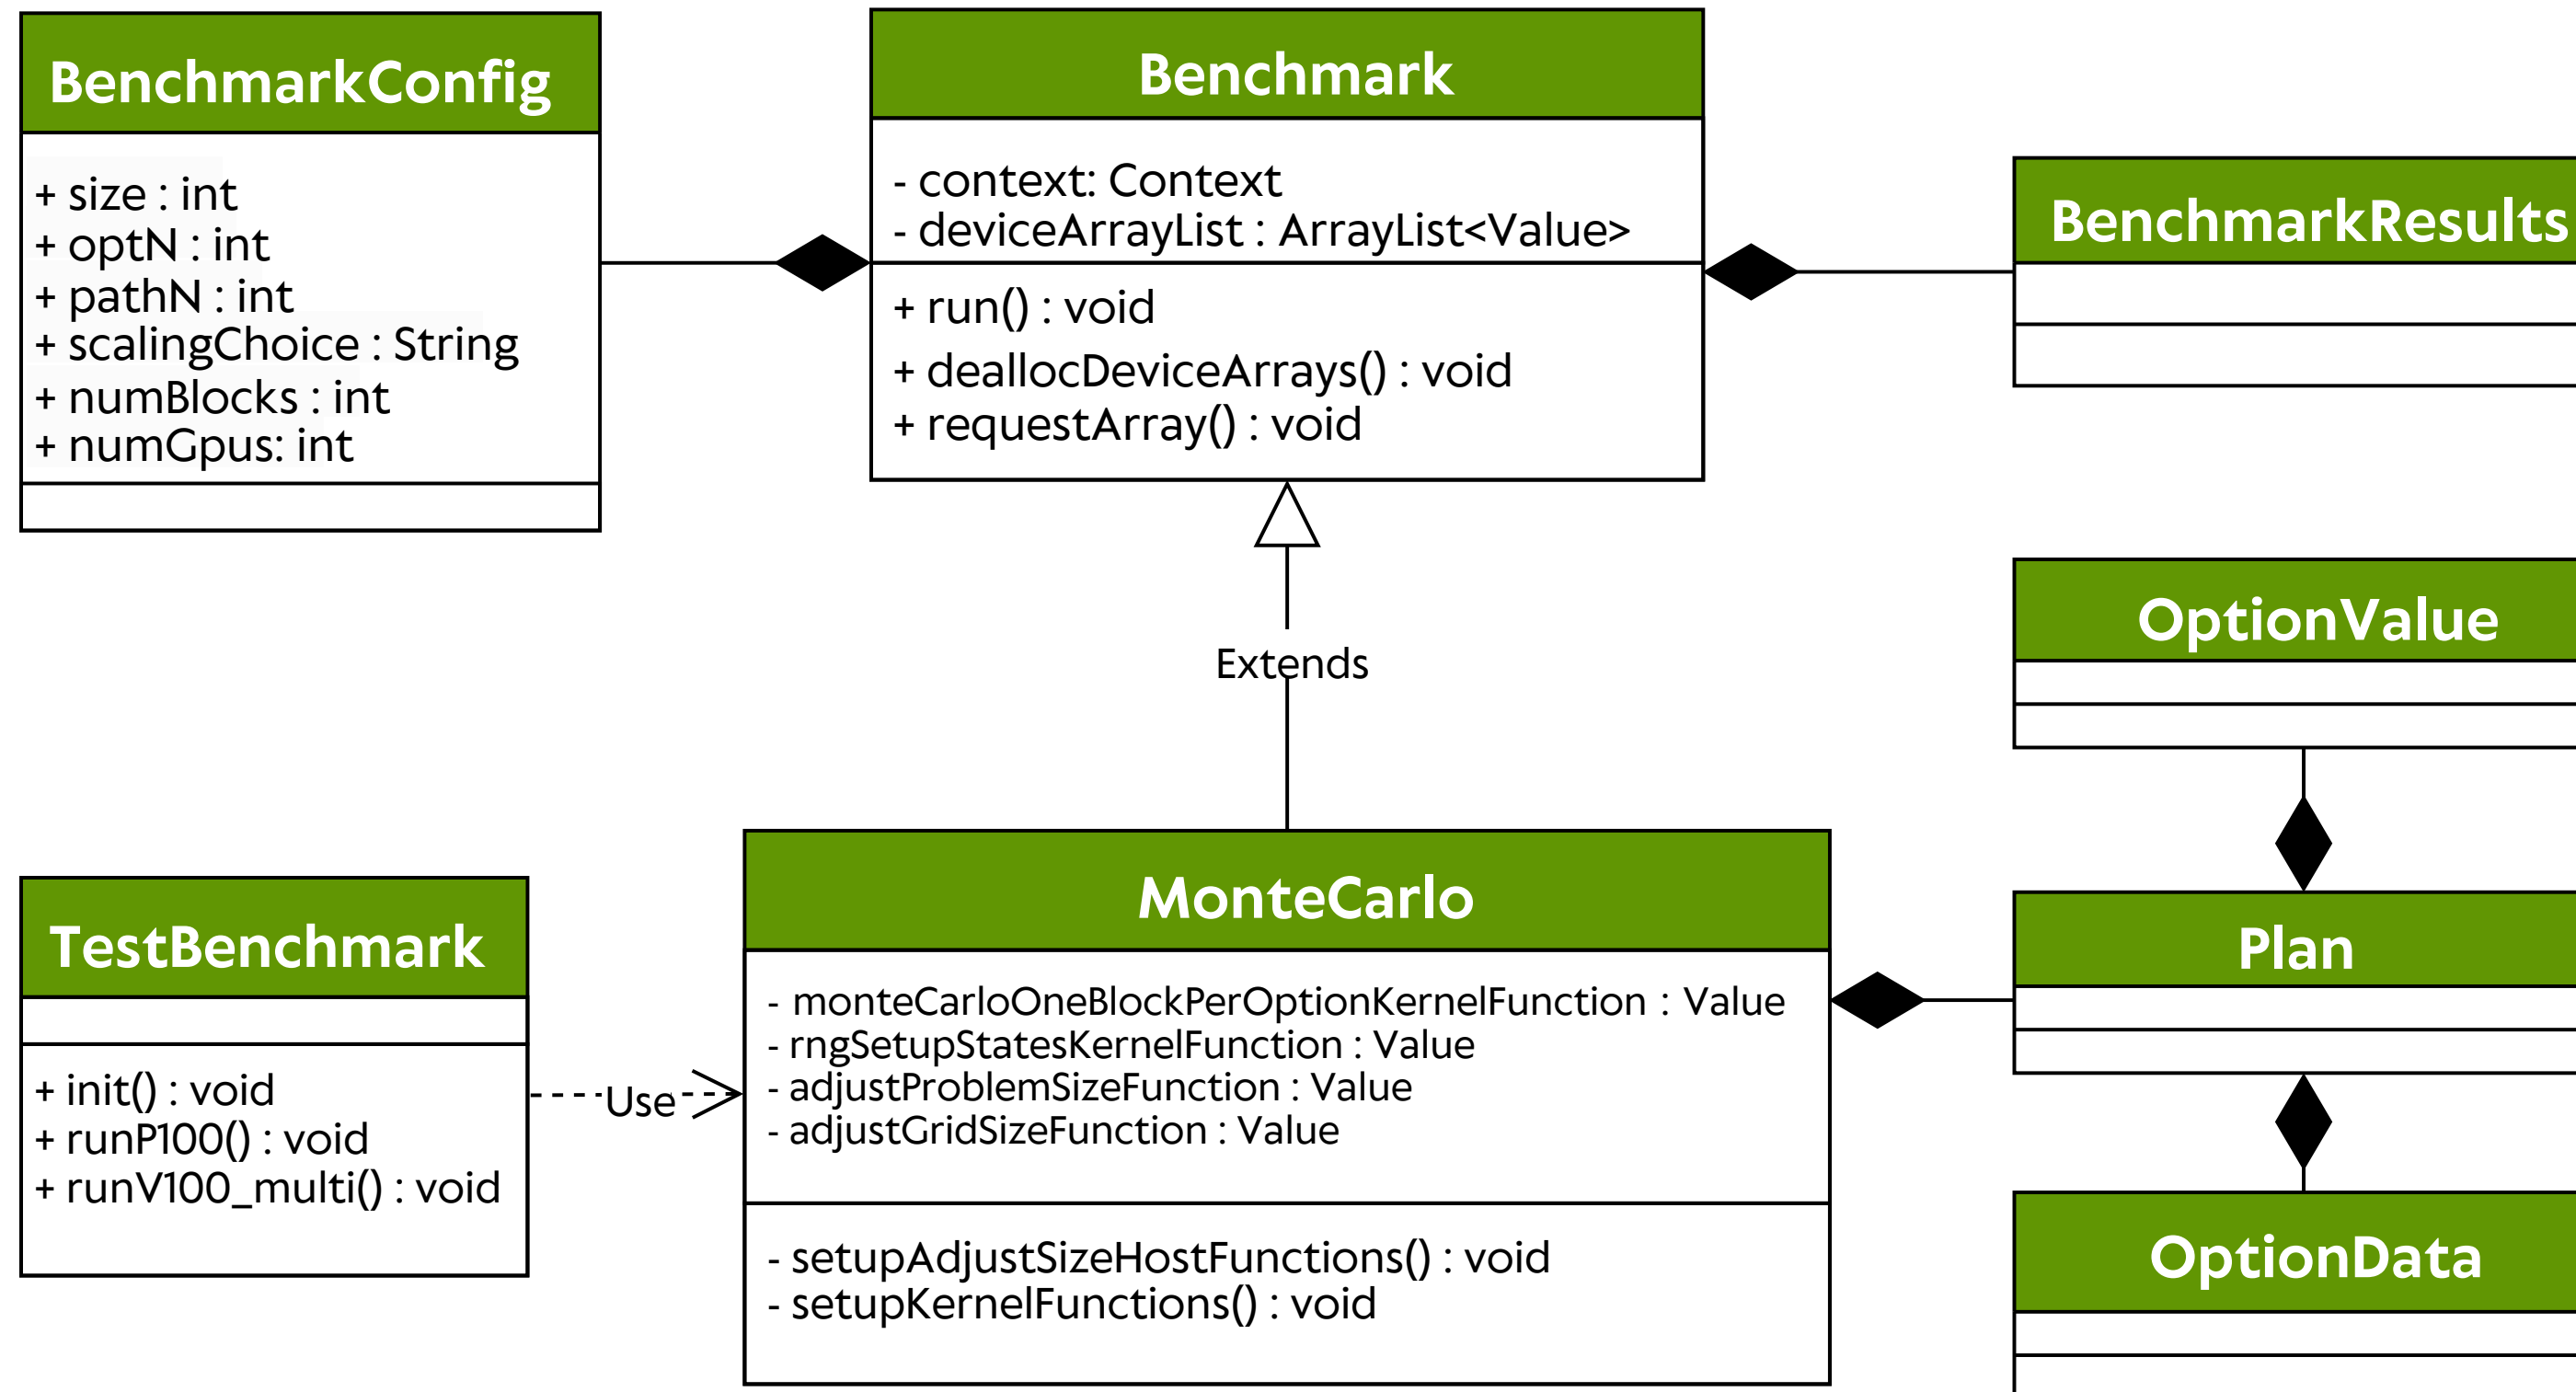
GrCUDA

Implementation

From the C++ implementation to
Java

- ◆ Abstraction of CUDA code to other languages
- ◆ Unified Memory automatically managed by the underlying system
- ◆ Memory regions stored into DeviceArrays

Our Benchmark Design



Best Practice

Binding kernels and functions
from files

Binding from files


CUDA kernels are written separately, compiled into a separate file, and bound using a binding function.

Intuition


Data structures and functions of external libraries are not supported by the standard building from Java strings.

Binding from files is used to overcome library import constraints.

Results - Comparing GPUs Number

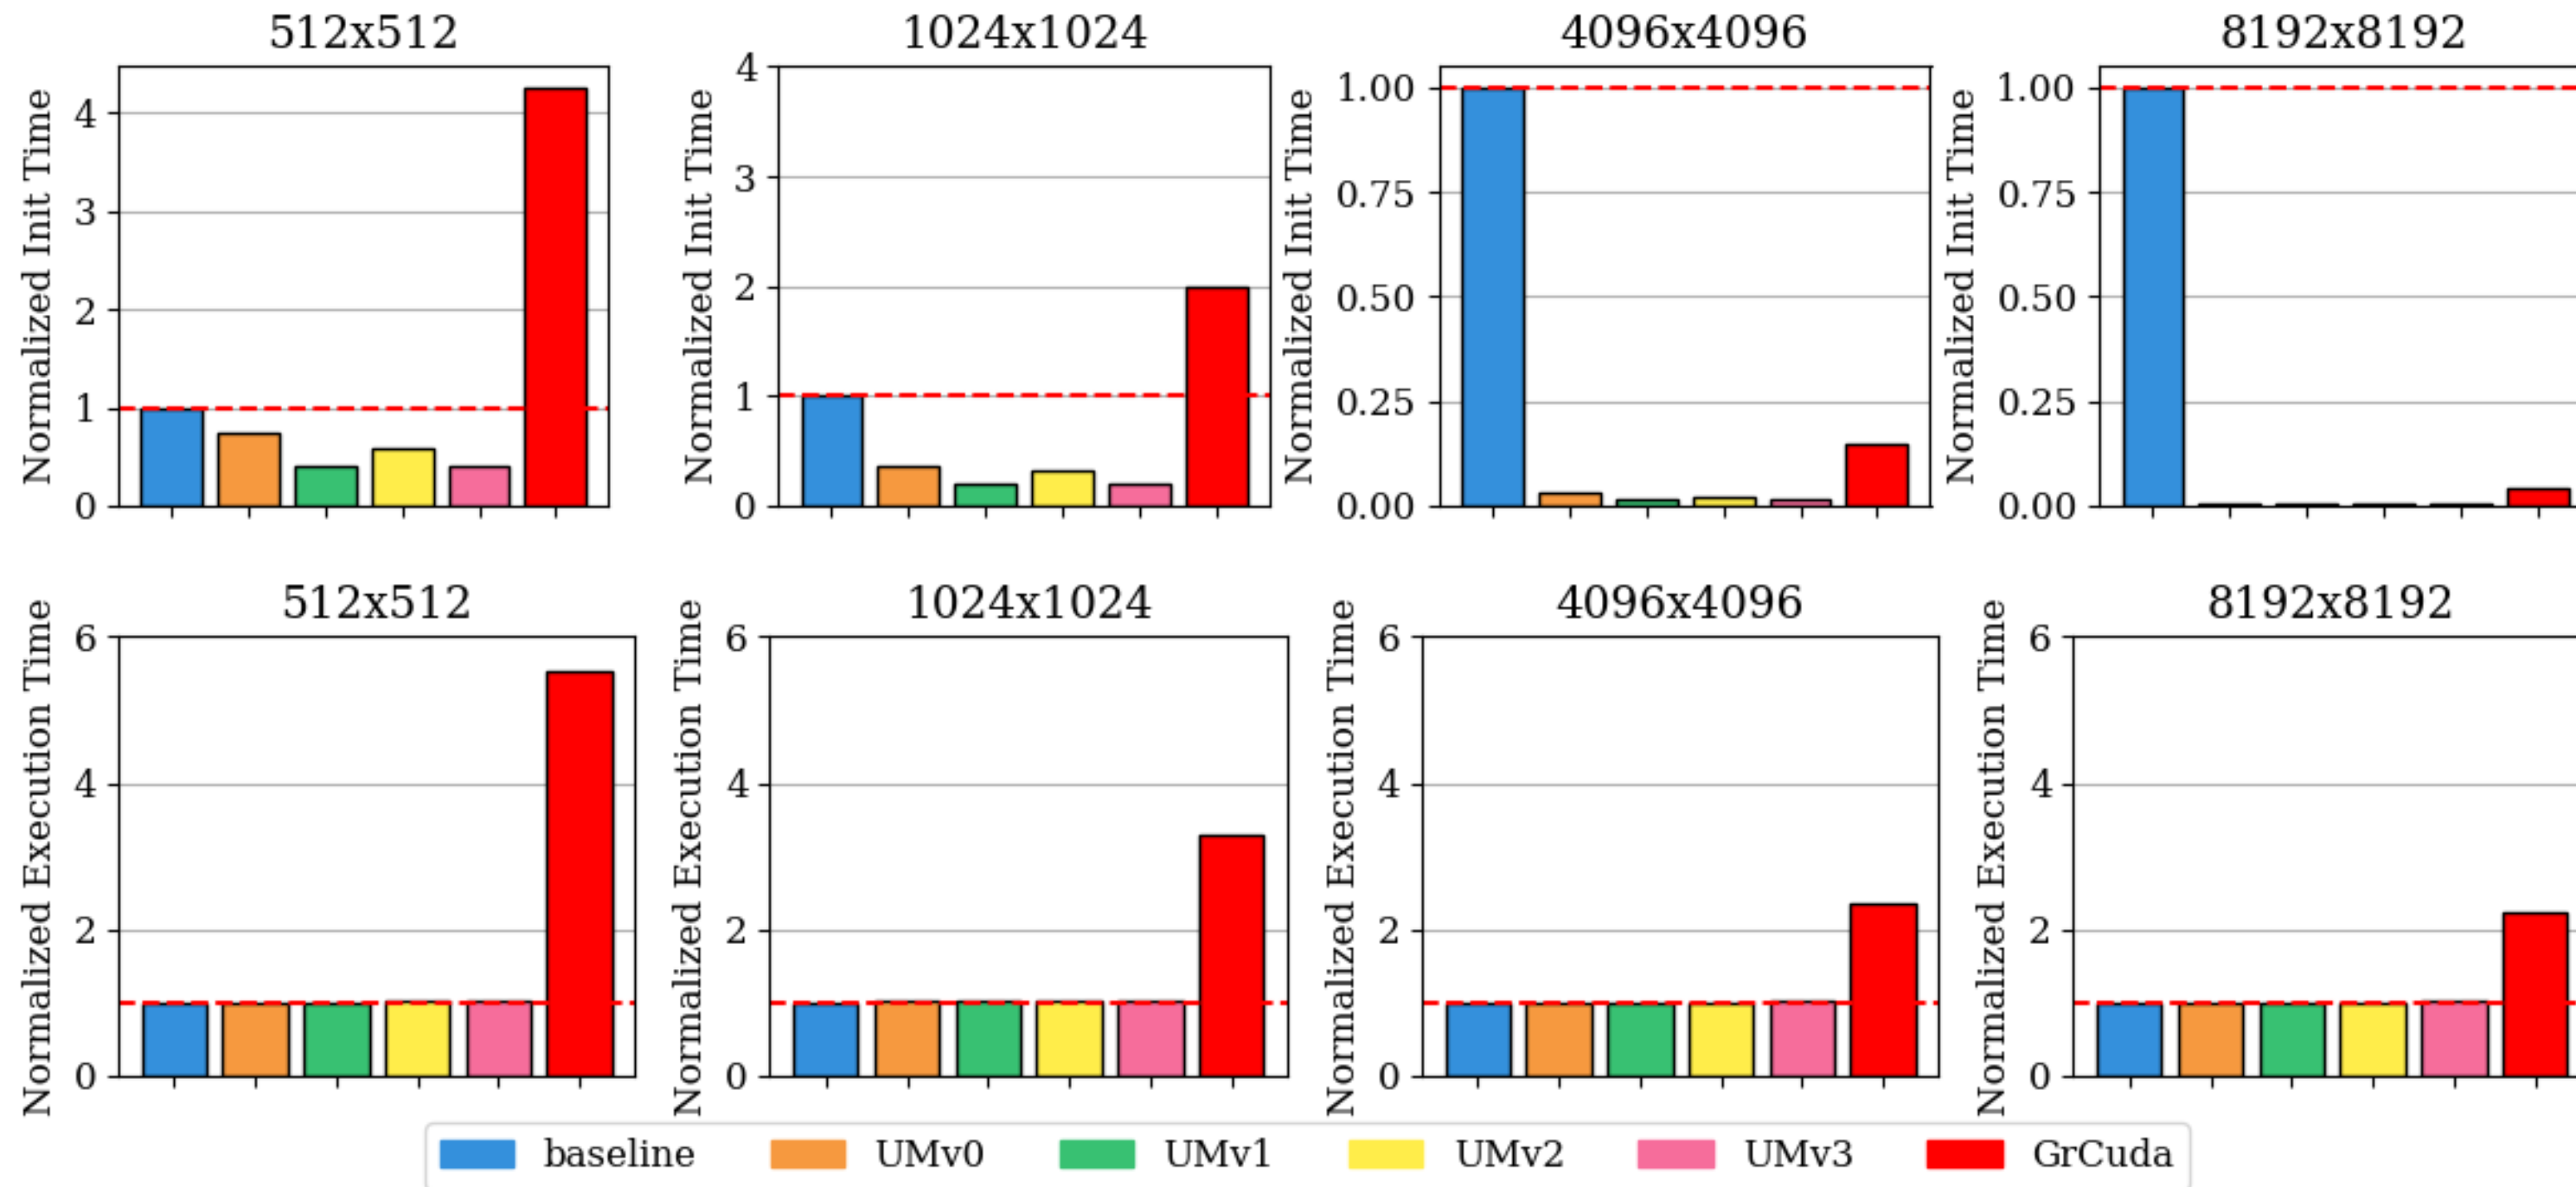


1.5x - 3x
Increased GrCuda
Execution Time



As GPU number increases
GrCuda performance gets worse

Results - Comparing Sizes



Performance
Improvement
as size increases

Streamed Strong Scaling

Our Conclusions

#1

Unified Memory
optimizations reduced
Initialization Time

#2

Unified Memory
optimizations introduction
didn't alter execution time.

#3

Benchmarks execution
patterns matter when
using Unified Memory

#4

GrCuda had worse
performances but
improved with size

Thank **you** for your **attention**!



Andrea Paparella
andrea1.paparella@polimi.mail.it



Andrea Piras
andrea6.piras@polimi.mail.it